

Name: **Krishna Gadam**

Div: **BE9-Q9**

Roll no: **43124**

Title: **Assignment 2: Implementing Feedforward neural networks with Keras and TensorFlow**

```
In [8]: #installations
from sklearn.preprocessing import LabelBinarizer
from sklearn.metrics import classification_report
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.optimizers import SGD
from tensorflow.keras.datasets import mnist
from tensorflow.keras import backend as K
import matplotlib.pyplot as plt
import numpy as np
```

```
In [9]: #grabbing the mnist dataset
((X_train, Y_train), (X_test, Y_test)) = mnist.load_data()
X_train = X_train.reshape((X_train.shape[0], 28 * 28 * 1))
X_test = X_test.reshape((X_test.shape[0], 28 * 28 * 1))
X_train = X_train.astype("float32") / 255.0
X_test = X_test.astype("float32") / 255.0
```

```
In [10]: lb = LabelBinarizer()
Y_train = lb.fit_transform(Y_train)
Y_test = lb.transform(Y_test)
```

```
In [11]: #building the model
model = Sequential()
model.add(Dense(128, input_shape=(784,), activation="sigmoid"))
model.add(Dense(64, activation="sigmoid"))
model.add(Dense(10, activation="softmax"))
```

```
In [15]: sgd = SGD(0.01)
epochs=10
model.compile(loss="categorical_crossentropy", optimizer=sgd,metrics=["accuracy"])
H = model.fit(X_train, Y_train, validation_data=(X_test, Y_test),epochs=epochs, t
```

Epoch 1/10

469/469 [=====] - 3s 6ms/step - loss: 1.7774 - accuracy: 0.5913 - val_loss: 1.6768 - val_accuracy: 0.6130

Epoch 2/10

469/469 [=====] - 2s 5ms/step - loss: 1.5965 - accuracy: 0.6282 - val_loss: 1.4964 - val_accuracy: 0.6528

Epoch 3/10

469/469 [=====] - 2s 5ms/step - loss: 1.4255 - accuracy: 0.6666 - val_loss: 1.3359 - val_accuracy: 0.6909

Epoch 4/10

469/469 [=====] - 2s 5ms/step - loss: 1.2772 - accuracy: 0.7010 - val_loss: 1.2009 - val_accuracy: 0.7123

Epoch 5/10

469/469 [=====] - 2s 5ms/step - loss: 1.1540 - accuracy: 0.7259 - val_loss: 1.0890 - val_accuracy: 0.7391

Epoch 6/10

469/469 [=====] - 3s 5ms/step - loss: 1.0531 - accuracy: 0.7453 - val_loss: 0.9984 - val_accuracy: 0.7653

Epoch 7/10

469/469 [=====] - 2s 5ms/step - loss: 0.9705 - accuracy: 0.7646 - val_loss: 0.9232 - val_accuracy: 0.7744

Epoch 8/10

469/469 [=====] - 2s 5ms/step - loss: 0.9016 - accuracy: 0.7781 - val_loss: 0.8599 - val_accuracy: 0.7926

Epoch 9/10

469/469 [=====] - 2s 5ms/step - loss: 0.8431 - accuracy: 0.7917 - val_loss: 0.8062 - val_accuracy: 0.8018

Epoch 10/10

469/469 [=====] - 2s 5ms/step - loss: 0.7926 - accuracy: 0.8022 - val_loss: 0.7584 - val_accuracy: 0.8154

```
In [16]: #making the predictions
predictions = model.predict(X_test, batch_size=128)
print(classification_report(Y_test.argmax(axis=1),predictions.argmax(axis=1),target_names=
```

```
79/79 [=====] - 0s 3ms/step
```

	precision	recall	f1-score	support
0	0.86	0.96	0.91	980
1	0.87	0.98	0.92	1135
2	0.85	0.77	0.81	1032
3	0.76	0.81	0.79	1010
4	0.73	0.84	0.78	982
5	0.81	0.61	0.70	892
6	0.87	0.90	0.89	958
7	0.84	0.86	0.85	1028
8	0.78	0.71	0.74	974
9	0.76	0.67	0.71	1009
accuracy			0.82	10000
macro avg	0.81	0.81	0.81	10000
weighted avg	0.81	0.82	0.81	10000

```
In [17]: #plotting the training loss and accuracy
plt.style.use("ggplot")
plt.figure()
plt.plot(np.arange(0, epochs), H.history["loss"], label="train_loss")
plt.plot(np.arange(0, epochs), H.history["val_loss"], label="val_loss")
plt.plot(np.arange(0, epochs), H.history["accuracy"], label="train_acc")
plt.plot(np.arange(0, epochs), H.history["val_accuracy"], label="val_acc")
plt.title("Training Loss and Accuracy")
plt.xlabel("Epoch #")
plt.ylabel("Loss/Accuracy")
plt.legend()
```

Out[17]: <matplotlib.legend.Legend at 0x7f55a707da10>

